

Beyond Boolean Intersections: A Type-Safe Verified Bisimulation Between D-Separation Trails and Moral-Graph Reachability

ANONYMOUS AUTHOR(S)

1 Introduction

Consider the simplest possible directed acyclic graph: a chain of three nodes $0 \rightarrow 1 \rightarrow 2$. Let $X = \{0\}$, $Y = \{1\}$, and $Z = \{0\}$. Does Z d-separate X from Y ?

The two standard textbooks give different answers:

- **Trail blocking** [2] says: examine every undirected trail from X to Y ; if every internal triple is blocked by Z , the sets are d-separated. The one-edge trail $0 \rightarrow 1$ has no internal triple, so it is not blocked. Hence X and Y are *not* d-separated.
- **Moral graph separation** [1] says: take the ancestral subgraph of $X \cup Y \cup Z$, moralize it, delete Z , and check whether X and Y are disconnected. Node 0 is in Z , so it is deleted. The remaining graph has no path from 0 to 1, so X and Y *are* d-separated.

We mechanize this discrepancy and its repair, treating the two characterizations not as interchangeable definitions but as a *verified bisimulation* between two semantic layers:

- an **operational semantics** of active trails (the “execution” layer of information flow);
- a **static analysis** of moralized-graph reachability (the “compiled” layer).

The forward direction (soundness) is a *certified trace optimizer* that compiles an operational trace into static reachability evidence. The reverse direction (completeness) is a *constructive decompiler of semantic flow*: it reconstructs a concrete execution witness—an exploitability certificate for the information leak—from static reachability alone. The counterexample says exactly that the two languages are *not* bisimilar without a well-typed domain restriction.

1.1 Contributions

Our contributions are:

- (1) We formalize both d-separation characterizations in Lean 4 over finite DAGs with natural-number nodes, organized as a three-layer stack: type-safe graph syntax, certified intermediate representations, and constructive witness-extraction pipelines.
- (2) We construct and machine-check a concrete counterexample, proving that unrestricted equivalence is false and identifying the exact repair ($\text{DisjointSets } X \ Y \ Z$).
- (3) We mechanize the **forward bisimulation** (soundness): an explicit compilation pipeline $\text{Trail} \rightarrow \text{BayesBallPath} \rightarrow \text{MAGWalk} \rightarrow \text{Reachable}$, certified by node-survival proofs.
- (4) We mechanize the **reverse bisimulation** (witness-producing completeness): a decompilation pipeline $\text{Reachable} \rightarrow \text{StaticRoute} \rightarrow \text{OpenTrace} \rightarrow \text{ActiveRoute} \rightarrow \exists \text{ Trail}, \neg \text{isBlocked}$. Every stage is type-indexed and constructive, including a proved bad-collider normalization (`exists_split`, `ancestor_escape`, `route_improves_of_bad`) that strictly decreases a well-founded measure. The reverse pipeline carries no **sorry**.
- (5) We close the **full equivalence** (`dSeparated_iff_dSeparates`): on the pairwise-disjoint domain, moral-graph separation holds if and only if every trail is blocked, with both directions discharged by the optimizer and the decompiler.

- (6) We delimit the boundary to **quantitative information flow**: the bridge from d-separation to probabilistic conditional independence is given as an explicit Lean scaffold (`dSeparation_implies_conc`) and is deliberately left undischarged, marking precisely where graph semantics ends and measure-theoretic work begins.

All results are available as a standalone Lean 4 package. The decompiler produces not a mere existence proof but a computational witness—a Σ -type containing the actual trail and the proof that it is unblocked.

2 A Three-Layer Architecture

Our formal development is organized in three layers that mirror the architecture of a verified compiler.

2.1 Layer 1: Type-Safe Graph Syntax

Nodes are natural numbers, edges are a finite set of ordered pairs ($\text{Finset } (\mathbb{N} \times \mathbb{N})$), and acyclicity is well-foundedness of the edge relation. The key safety invariant is $\text{DisjointSets } X \ Y \ Z$, which requires pairwise disjointness of the three node sets. In the language of substructural types, this is an *affine ownership constraint*: the conditioned variables Z are borrowed by the query and cannot alias the source X or sink Y .

Listing 1. DAG structure and safety predicate

```

structure DAG where
  nodes : Finset \N
  edges : Finset (\N \times \N)
  edges_subset : \forall {u v}, (u,v) \in edges -> u \in nodes \wedge v \in nodes
  acyclic : WellFounded fun u v => (u,v) \in edges

def DisjointSets (X Y Z : Finset \N) : Prop :=
  Disjoint X Y \wedge Disjoint X Z \wedge Disjoint Y Z

```

2.2 Layer 2: Certified Intermediate Representations

Between operational trails and graph reachability sits a family of IRs that preserve just enough structure for certified translation.

MAGWalk. A compressed walk language equivalent to moral-graph reachability. It has three constructors: reflexivity, a single directed or reversed edge (`MAGWalk.single`), and a “moral jump” over an active collider (`MAGWalk.jump`). The collider node is skipped, which is why compression is correct.

StaticRoute. A type-valued decompiler IR that keeps the witness hidden by graph reachability. Each step is either a forward edge, a backward edge, or a moral jump storing the common child:

Listing 2. Static route IR (excerpt)

```


inductive StaticStep (G : DAG) (X Y Z : Finset \N) : \N -> \N -> Type where
  | directForward (hEdge : G.HasEdge u v) ...
  | directBackward (hEdge : G.HasEdge v u) ...
  | moralJump (huw : G.HasEdge u w) (hvw : G.HasEdge v w)
    (hw : w \in G.ancestralSubgraphNodes (X \cup Y \cup Z)) ...

```

OpenTrace. The compiler target before converting to an active route. Each cons node stores its oriented DAG traversal *and* a local non-blocking proof, making *OpenTrace* a type-valued certificate that every junction is *Z*-open.

2.3 Layer 3: Witness-Extraction Pipelines

The two pipelines are summarized in Figure 1.



Reverse: Decomilation of Semantic Flow

Fig. 1. The two witness-extraction pipelines. Solid arrows: forward (soundness). Dashed arrows: reverse (completeness). Every arrow, including the *normalize* step, is a total constructive function in Lean; the *normalize* step is a proved bad-collider minimization (Section 6).

3 Background: The Two Criteria

3.1 Trail Blocking (Operational Semantics)

An undirected *trail* from x to y in a DAG G is a finite walk that may follow edges forward or backward. A consecutive triple $a \rightarrow b \rightarrow c$ or $a \leftarrow b \rightarrow c$ is a *non-collider*; $a \rightarrow b \leftarrow c$ is a *collider*. A triple is *blocked* by Z if:

- it is a non-collider whose middle node is in Z ; or
- it is a collider whose middle node and all descendants are disjoint from Z .

A trail is *blocked* by Z if some internal triple is blocked. Then Z *d-separates* X from Y in the trail sense if every trail from X to Y is blocked by Z .

3.2 Moral Graph Separation (Static Analysis)


The *ancestral subgraph* of S is the subgraph induced by all ancestors of nodes in S . Its *moral graph* is obtained by (1) forgetting edge directions and (2) connecting all co-parents (pairs of nodes with a common child). The *d-separation graph* is the moral graph of the ancestral subgraph of $X \cup Y \cup Z$ with Z removed. Then Z *d-separates* X from Y in the graph sense if no node of X is connected to any node of Y in this graph.

3.3 Textbook Equivalence and Its Caveat

Standard references state that the two criteria are equivalent under the implicit assumption that X , Y , and Z are pairwise disjoint. In applications, however, the disjointness precondition is often omitted. Our counterexample (Section 4) shows that this omission is not benign.

4 The Endpoint Caveat: A Counterexample

Our counterexample lives on the simplest non-trivial DAG:



Let $X = \{0\}$, $Y = \{1\}$, $Z = \{0\}$.

THEOREM 4.1 (ENDPOINT-IN-Z COUNTEREXAMPLE). *There exists a DAG G and finite sets $X, Y, Z \subseteq \mathbb{N}$ such that G d -separates X from Y in the moral-graph sense, but Z does not d -separate X from Y in the trail-blocking sense.*

Listing 3. Mechanized counterexample

```

theorem dsep_complete_endpoint_in_Z_counterexample :
  \exists (G : DAG) (X Y Z : Finset \N),
    DAG.dSeparated G X Y Z \wedge \neg dSeparates G X Y Z := by
  refine <chain3, {0}, {1}, {0}, ?_, ?_>
  ...

```

Why the disagreement happens. Trail blocking only inspects *internal* triples. The one-edge trail $0 \rightarrow 1$ has no internal triple at all, so nothing can block it. By contrast, the moral-graph construction deletes Z unconditionally; since $0 \in Z$, the edge $(0, 1)$ disappears from the moralized graph, and 0 and 1 are disconnected.

The standard repair. If we require $X \cap Z = \emptyset$ and $Y \cap Z = \emptyset$, the start node of any trail is never deleted, and the endpoint caveat disappears.

5 Forward Bisimulation: The Certified Trace Optimizer

Under DisjointSets $X Y Z$, we mechanize the soundness direction:

$$\text{DAG.dSeparated}(G, X, Y, Z) \implies \text{dSeparates}(G, X, Y, Z).$$

The proof follows a three-stage compilation pipeline.

Stage 1: Trail $\rightarrow$ Bayes-ball path. If a trail t from $x \in X$ to $y \in Y$ is not blocked by Z , and $x \notin Z$ (guaranteed by Disjoint $X Z$), then we lift t to an explicit Bayes-ball state path. The construction is by induction on the trail, analyzing each directional triple to show that the Bayes-ball state machine can traverse it.

Listing 4. Trail to Bayes-ball path

```

def bayesBallPath_of_active_trail_outOf
  (t : Trail G u v) (h_active : \neg t.isBlocked Z)
  (huZ : u \notin Z) :
  \Sigma final_dir, BayesBallPath G Z (u, outOf) (v, final_dir)

```

Stage 2: Certified membership. The Bayes-ball path contains intermediate states that must “survive” deletion of Z . We prove a certified version that guarantees every RequiredState node belongs to the d -separation graph nodes:

Listing 5. Certified Bayes-ball path

```

def bayesBallPathCert_of_active_trail_outOf
  (t : Trail G u v) (h_active : \neg t.isBlocked Z)
  (huZ : u \notin Z)

```

```
(huA : u \in G.ancestralSubgraphNodes (X \cup Y \cup Z))
(hvD : v \in G.dSeparationGraphNodes X Y Z) :
\Sigma final_dir, {p : BayesBallPath ... //
  \forall q, BayesBallPath.RequiredState p q ->
    q.1 \in G.dSeparationGraphNodes X Y Z}
```

Stage 3: Compress to MAGWalk. A Bayes-ball path is a sequence of single-edge steps. The compress algorithm scans it with a two-step window: ordinary non-collider windows become a MAGWalk. single, while active-collider windows $(a, _) \rightarrow (b, \text{into}) \rightarrow (c, \text{outOf})$ become a single MAGWalk. jump over the collider b . The collider itself need not survive deletion of Z , which is why this compression is correct.

THEOREM 5.1 (SOUNDNESS). *For any DAG G and pairwise-disjoint finite sets $X, Y, Z \subseteq \mathbb{N}$, if G d -separates X from Y in the moral-graph sense, then every trail from X to Y is blocked by Z .*

Listing 6. Top-level soundness theorem

```
theorem dSeparated_of_dSeparated_disjoint
(hXYZ : DisjointSets X Y Z)
(hsep : DAG.dSeparated G X Y Z) : dSeparates G X Y Z
```

6 Reverse Bisimulation: Decomilation of Semantic Flow

The reverse direction is the decompiler: from $\neg \text{DAG.dSeparated}$ it constructively reconstructs an active-trail witness—an exploitability certificate showing that the conditioning set Z fails to block the flow. The pipeline is:

Reachable $\rightarrow$ StaticRoute $\rightarrow$ NormalizedStaticRoute $\rightarrow$ OpenTrace $\rightarrow$ ActiveRoute $\rightarrow \exists \text{Trail}, \neg \text{isBlocked}$

6.1 From Reachability to StaticRoute

Graph reachability is a **Prop**-valued relation that hides the structure of each step. We introduce StaticStep and StaticRoute as type-valued evidence: a StaticRoute is an explicit list of directed edges, reversed edges, and moral jumps, each carrying the witness needed for later stages.

Listing 7. Reachability to static route

```
theorem nonemptyStaticRoute_of_dSeparationGraph_reachable
(h : (G.dSeparationGraph X Y Z).Reachable u v) :
Nonempty (StaticRoute G X Y Z u v)
```

6.2 Normalization: Eliminating Bad Colliders

A “bad collider” is a collider window whose middle node and all descendants are disjoint from Z . Such a collider is blocked, so it must be eliminated. Normalization may change endpoints: if a bad collider leaks ancestry into X or Y , we reroute the prefix or suffix to a different node in that set.

We package static routes as endpoint-carrying witnesses (StaticRouteWitness) and define a bad-collider count routeBadCount. The minimization is a Nat . find wrapper: if every nonzero-count witness can be strictly improved, a zero-count witness exists.

Listing 8. Bad-collider minimization wrapper

```
theorem normalized_route_exists_of_improves
(hwit : \exists w : StaticRouteWitness G X Y Z, True)
(himprove : \forall w, routeBadCount w \neq 0 ->
```

```

\exists w', routeBadCount w' < routeBadCount w) :
\exists w, routeBadCount w = 0

```

The improvement obligation is discharged by the core normalization theorem, which is fully proved (no **sorry**):

Listing 9. Strict improvement for non-normalized routes

```

theorem route_improves_of_bad
  (w : StaticRouteWitness G X Y Z) (hbad : routeBadCount w \neq 0) :
  \exists w' : StaticRouteWitness G X Y Z,
    routeBadCount w' < routeBadCount w

```

The proof is a short dispatcher over three established lemmas. `exists_split` extracts the first bad collider as a `Split` carrying a zero-bad prefix and a count bound. `ancestor_escape` shows the bad child’s leaked ancestry must reach a node in X or Y (its descendant cone misses Z , so it cannot dead-end). `bad_child_survives` and `escape_path_survives` prove that the rerouted escape chain stays inside the d-separation graph nodes. Reattaching the escape chain—backward into X or forward into Y —removes a moral jump of cost 1 and adds chains of cost 0; the strict decrease then follows by `countBadColliders_backwardEscape_append_suffix_lt` (resp. `countBadColliders_prefix_append_forwardEscape`) and `omega`. As a deliberate design choice, the witness endpoints may change: rerouting is allowed to land on a *different* node of X or Y , which is what makes the escape construction total.

6.3 From Zero-Bad StaticRoute to OpenTrace to ActiveRoute

A static route with zero bad colliders compiles directly to an `OpenTrace`, because every junction is locally open:

Listing 10. Zero-bad compiler

```

theorem openTrace_of_countBadColliders_zero
  (hzero : countBadColliders arrival route = 0) :
  Nonempty (\Sigma finalDir, OpenTrace G Z (x, arrival) (y, finalDir))

```

The `OpenTrace` is then converted to a `BayesBallPath` and wrapped as an `ActiveRoute`. Finally, `ActiveRoute.to_activeTrail` yields the existential witness:

Listing 11. Final decompiler theorem

```

theorem activeWitness_of_not_dSeparated
  (hnot : \neg DAG.dSeparated G X Y Z) : ActiveWitness G X Y Z

```

The result is a constructive Σ -type containing the actual nodes $x \in X$, $y \in Y$, the final direction, and the active route.

6.4 The Closed Equivalence

Composing the optimizer (Section 5) and the decompiler yields the full equivalence on the well-typed domain—the bisimulation statement: on pairwise-disjoint X, Y, Z , the operational and static languages agree exactly.

Listing 12. Full d-separation equivalence (no proof debt)

```

theorem dSeparated_iff_dSeparates
  (hXYZ : DisjointSets X Y Z) :
  DAG.dSeparated G X Y Z <-> dSeparates G X Y Z

```

[Verification status] The full Lean development typechecks (lake build). The graph-semantic core—both directions of `dSeparated_iff_dSeparates` and the entire reverse synthesis stack—contains no `sorry`. The only `sorry`s in the package are two *intentional* scaffold stubs in `InfoTheoryBridge.lean` (Section 8), which are out of scope for the equivalence claims.

7 The PL Dictionary: Reframing Causal Graphs

To situate this work for a POPL audience, we reframe traditional causal-graph vocabulary in programming-languages terms:

Causal Inference	PL Framing
DAG	Dataflow AST / Dependency Graph Syntax
<code>DisjointSets X Y Z</code>	Affine ownership constraint (no aliasing of Z)
Active Trail	Operational trace of semantics
MAGWalk	Reachability abstract IR (compressed moral-graph walk)
Collider Jump (<code>MAGWalk.jump</code>)	Peephole optimization / semantic fold
<code>Trail → BayesBallPath</code>	Typestate tracking
<code>BayesBallPath → MAGWalk</code>	Certified trace optimizer
<code>StaticRoute → OpenTrace → ActiveRoute</code>	Semantic-flow decompiler / exploitability-certificate generator
Cut-Set / Mutual Information	Quantitative Information Flow (QIF) capacity bounds

Table 1. Reframing causal-graph concepts for a POPL audience.

8 Related Work and Future Directions

Causal inference formalization. Causal reasoning has been formalized in Isabelle/HOL, Coq, and recently in Lean 4. To our knowledge, no prior formalization treats the endpoint caveat, gives a machine-checked counterexample to unrestricted equivalence, or provides a constructive reverse decompiler.

Graph theory in proof assistants. Mathlib 4 contains a mature `SimpleGraph` library. Our DAG definition is more algorithmic (finite node/edge sets, well-founded acyclicity) than the algebraic approach typical of Mathlib, making it convenient for constructive computation and witness extraction.

Quantitative Information Flow (QIF). Security analysis has moved from non-interference (“0 or 1”) to measuring leak in bits. The verified bisimulation here is intended as the *graph-semantic substrate* for such bounds: a closed d-separation equivalence is what licenses replacing a conditional-independence assumption by a checked graph property in a cut-set capacity argument. Crossing from graph reachability to a probabilistic leak bound requires the bridge

$$\text{d-separation} \Rightarrow \text{conditional independence} \Rightarrow \text{conditional DPI} \Rightarrow \text{cut-set bound}.$$

We state the first step of this bridge as an explicit Lean scaffold, `dSeparation_implies_conditional_independence` in `InfoTheoryBridge.lean`, together with stubs for `FinitePMF` and `MarkovCompatible`. This scaffold is deliberately *not* discharged in this artifact: it carries two intentional `sorry`s and is the sole proof debt in the package. The finite-measure and information-theory machinery (entropy, conditional mutual information, conditional DPI, dual/KKT certificates) lives in a separate verification stack; integrating it with the DAG semantics formalized here—unifying the two DAG definitions and proving the conditional-independence step—is the principal item of future work.

9 Conclusion

We formalized the two standard d-separation criteria in Lean 4 as a verified bisimulation between an operational semantics of active trails and a static analysis of moralized-graph reachability. We machine-checked a concrete counterexample to unrestricted equivalence, characterized the repair (pairwise disjointness as an affine ownership constraint), and discharged both directions: a certified trace optimizer for soundness and a constructive, witness-producing decompiler—with a proved bad-collider normalization—for completeness. The two compose into the closed theorem `dSeparated_iff_dSeparates`, and the graph-semantic core carries no proof debt. What remains is the deliberately delimited bridge to probabilistic conditional independence, stated here as a scaffold. Closing it would turn this graph-semantic substrate into machine-checked quantitative information-flow bounds.

References

- [1] Steffen L. Lauritzen. *Graphical Models*. Oxford University Press, 1996.
- [2] Judea Pearl. *Causality*. Cambridge University Press, 2nd edition, 2009.